



Project - Develop a RAG Application Using Knowledge Graph and Vector Search

Goal:

Many RAG applications are built today using vector search. Still, the latest research shows that incorporating knowledge graphs can help improve the performance and quality of the Retrieval-Augmented Generation (RAG) application. You will develop an application that leverages Knowledge Graphs and vector search. This application will use 10,000 Wikipedia articles as inputs and create a knowledge graph and vector search database for these articles. Then you will take a query as input, retrieve relevant information from the structured knowledge base you created, and generate accurate and contextually appropriate responses.

Project Expectations:

1. Understanding RAG Architecture:

- Students are expected to understand the RAG model, including its components: retriever and generator.
- They should learn how Knowledge Graphs can be used to enhance retrieval by providing structured, contextual information.

2. Implementing Vector Search:

- Implement a vector search mechanism using tools such as FAISS, Annoy, Pinecone, or other vector search libraries to retrieve semantically relevant information from unstructured data sources, in this case, Wikipedia articles
- Understand how to generate and index embeddings, chunking, and perform efficient similarity searches

3. Integration of Knowledge Graphs:

- Integrate a Knowledge Graph into the RAG pipeline. You will use DBpedia, YAGO, or other knowledge graphs (KG) of your choice, and use SPARQL to query the KG.
- You will incorporate retrieval using KG in your RAG pipeline.

4. Integration of LLM for generation of responses:

- Use a publicly available LLM to generate responses from the data retrieved for a given query. While you can use a paid LLM such as OpenAI's models, you are also free to use free models such as Microsoft Phi, Smaller version of LLaMA, or others

5. Application Development:

- Develop a functional application that uses the RAG framework to answer queries or generate content based on the retrieved data.
- The application should be able to handle a variety of inputs, demonstrating the flexibility and effectiveness of combining Knowledge Graphs and vector search in a RAG setup.

6. Evaluation and Optimization:



- Evaluate the performance of the RAG application in terms of accuracy, relevance, and response quality.
- Optimize both the retriever (vector search) and generator components to improve the overall application performance.

7. Documentation and Presentation:

- Maintain thorough documentation of the project, detailing the design decisions, implementation steps, challenges faced, and solutions developed.
- Prepare a presentation or demonstration of the final application, showcasing its capabilities and explaining the underlying technical details.

Steps to Achieve the Goal:

1. Initial Research and Learning:

- Study the RAG model architecture and its components (retriever and generator).
- Learn about Knowledge Graphs: their creation, querying (using SPARQL, for instance), and integration into NLP tasks.
- Understand vector search fundamentals and tools available for implementation.

2. Data Collection and Knowledge Graph Construction:

- Gather the dataset for your application, in this case Wikipedia articles
- Utilize an existing Knowledge Graph that can be used for retrieval purposes.

3. Implementation of Vector Search:

- Generate embeddings for data points using pre-trained models (e.g., BERT).
- Index the embeddings using a vector search library and implement the retrieval mechanism.

4. Integration of Components:

- Connect the retriever (vector search) and integrate the Knowledge Graph to enhance the retrieval process.
- Connect the output of the retriever with the generator (language model) to create structured responses. Ensure the generator has access to structured information.

5. Application Development:

- Implement the logic for query processing, retrieval, and response generation.
- Command line query input and piping the response to a text file is adequate for this project

6. Testing and Optimization:

- Test the application with various inputs to assess its performance.
- Refine the retrieval algorithms, tweak the Knowledge Graph integration, and adjust the generator model to optimize results.

7. Documentation and Final Presentation:

- Document the entire process, including code, methodologies, and findings.
- Prepare a final presentation showcasing the application, discussing the technical challenges, solutions, and potential future improvements.