

HACKEREARTH RAG APPLICATION – SPRINT 2

Ethan Villalovoz, Molly Iverson, Adam Shtrikman, Chandler Juego

Introduction

We will develop a RAG (Retrieval-Augmented Generation) application for HackerEarth that will utilize vector search, knowledge graphs, and a LLM to answer questions and generate content from a knowledge base of more than 10,000 Wikipedia articles.

Sprint Objectives

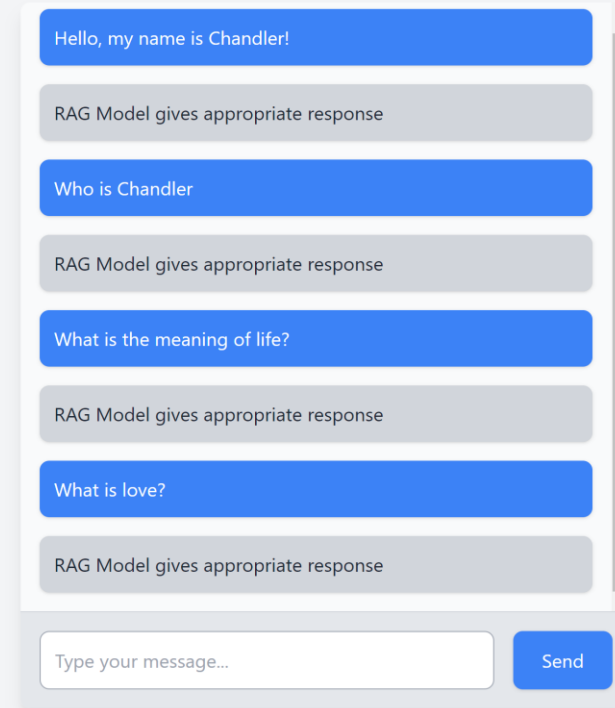
- **Embeddings Creation and Vector Search Preparation:** To implement efficient text embeddings, we researched and tested embedding methods, specifically focusing on BERT-based embeddings for use in the vector search module.
- **Frontend Prototype Development:** Developing a minimalistic chat-style interface to enable seamless interaction between users and the RAG application.
- **Knowledge Graph Querying:** Initial setup and testing of DBpedia integration to enable semantic data retrieval for more accurate responses.
- **System Architecture Refinement:** Documentation updates on the system architecture to ensure a clear breakdown of subsystems, including NLP, embeddings, and data processing workflows.

Feature Implementation

Team Member	Feature	Objective
Ethan	Embeddings	Embeddings Creation and Vector Search Preperation
Chandler	Frontend	Frontend Prototype Development
Molly	Knowledge Graph	Knowledge Graph Querying
Adam	Natural Language Processing	Query Processing

Figma Prototype & Client Feedback

- Minimalistic
chat/messaging system
with the RAG model



Kanban

Overview &

Contributions

Summary:

- **Molly**
 - Knowledge Graph code
 - Architecture diagram and subcomponent decomposition report sections
 - Unit and Integration testing report sections
 - Project Milestone document
- **Ethan**
 - Text embeddings research and code
 - Architecture introduction and system overview report sections
 - Testing strategy and environment requirements report sections
 - Sprint report
- **Chandler**
 - React Frontend prototype
 - User Interface design report section
 - System testing report section (functional, performance, and user acceptance testing)
 - Python code for extracting and embedding Wikipedia data
- **Adam**
 - Natural Language query processing code
 - Data Design and some subcomponent report sections
 - Testing overview, test objectives, and scope report sections
 - Python notebook with FAISS tutorial
- **All**
 - Meeting notes for client meetings

Ready 0 Estimate: 0

In progress 1 / 12 Estimate: 0

In review 0 / 12 Estimate: 0

Done 50 / 50 Estimate: 112

Write unit and integration tests section #54

Edit



Closed

mollyiverson/ACME10-HE-RAGApp Private

mollyiverson opened last week

edited by mollyiverson · Edits ·

As a developer

I need a clear strategy for unit and integration testing

So that I can ensure each component and their interactions function as expected

Details and Assumptions

- Unit testing will target individual handlers (LLM Handler, KG Handler, VS Handler, etc.), isolating them to verify accuracy in processing inputs and outputs.
- Integration testing will focus on testing pairs or groups of components, such as the KG Handler with the VS Handler, to confirm accurate data flow and compatibility.
- Dependencies for each component will be mocked during unit tests, while real connections will be tested in integration.

Acceptance Criteria

Given an isolated handler
When unit tests are executed
Then each handler's individual functionality is verified

Given two or more integrated handlers
When integration tests are conducted
Then data flows correctly between components, and expected outputs are observed without errors

Assignees

mollyiverson

Labels

documentation

Projects

Capstone Agile Planning

Status Done

Priority Choose an option

Size Choose an option

Estimate 3

Start date No date

End date No date

Sprint Sprint 2 • Oct 5 - 12

Milestone

No milestone

Development

Create a branch for this issue or link a pull request

This has been completed

ACME10-HE-RAGApp #56

Write System testing section (functional, performance, user acceptance)

3 Sprint 2

ACME10-HE-RAGApp #57

Create code for natural language processing for the query

3 Sprint 2

ACME10-HE-RAGApp #53

Write introduction for Testing and Acceptance plans report (overview, test objectives, scope)

3 Sprint 2

ACME10-HE-RAGApp #54

Write unit and integration tests section

3 Sprint 2

ACME10-HE-RAGApp #51

Integrate Knowledge Graphs in the RAG pipeline

3 Sprint 2

ACME10-HE-RAGApp #52

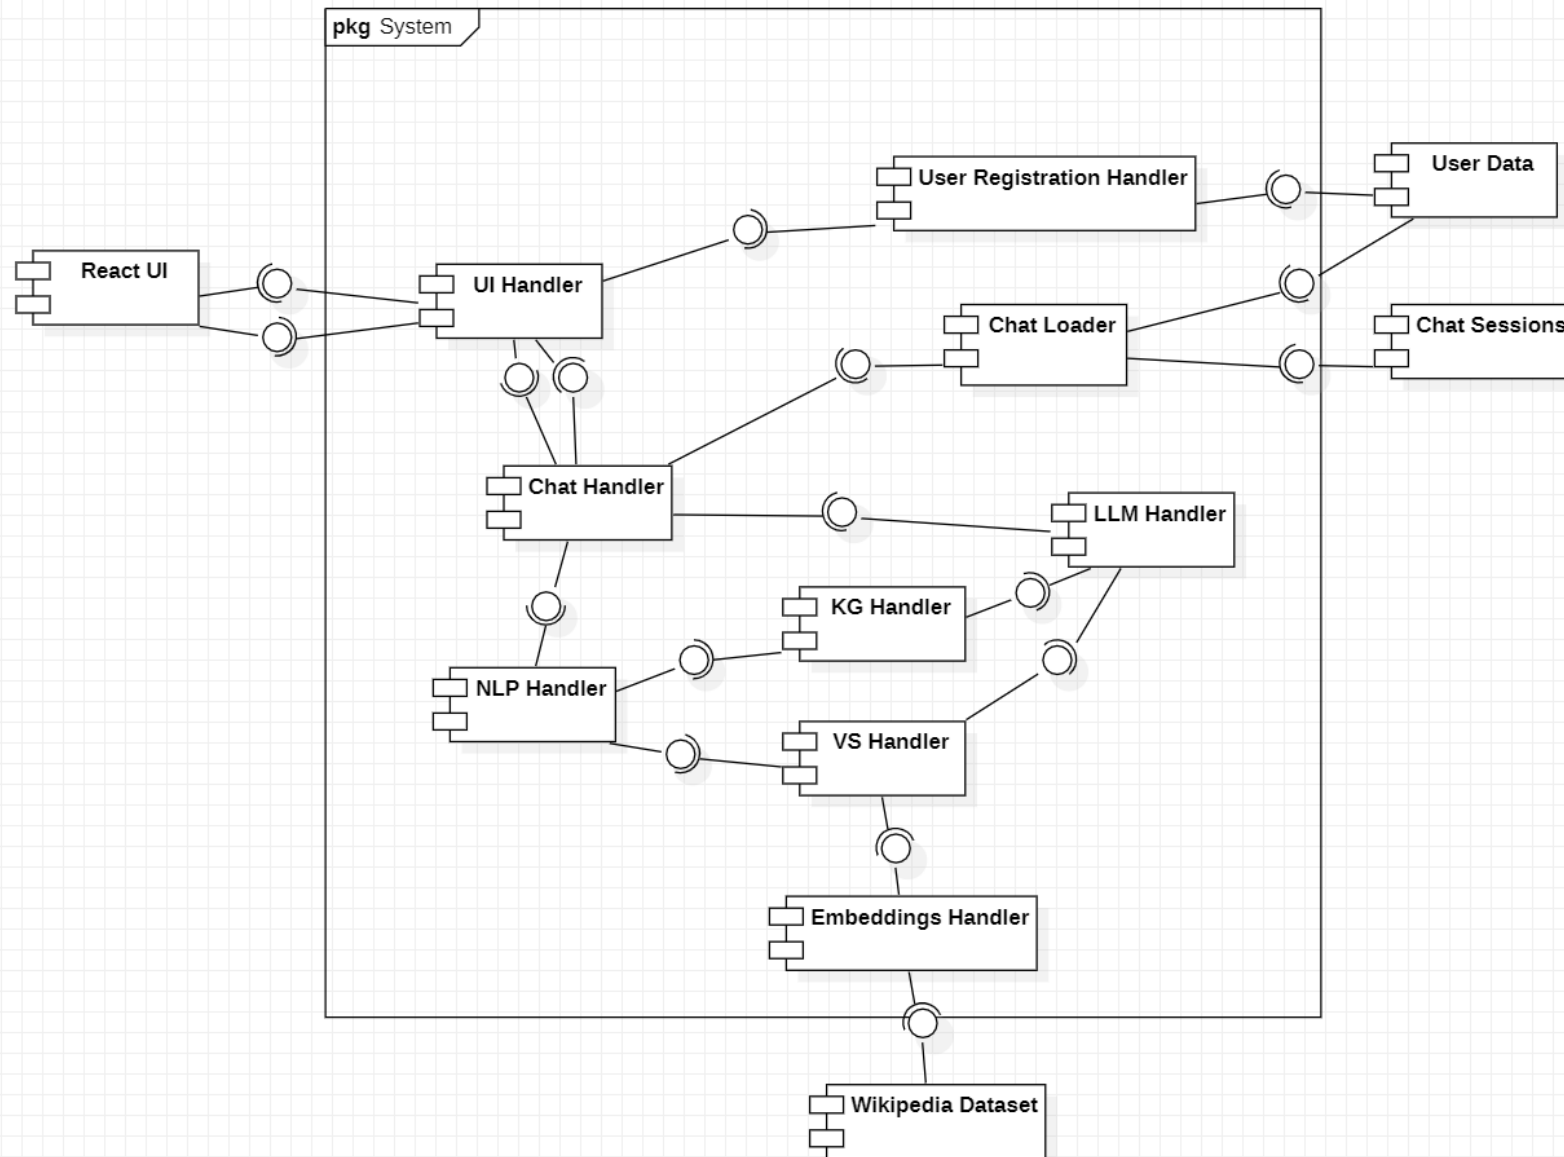
+ Add item

+ Add item

+ Add item

+ Add item

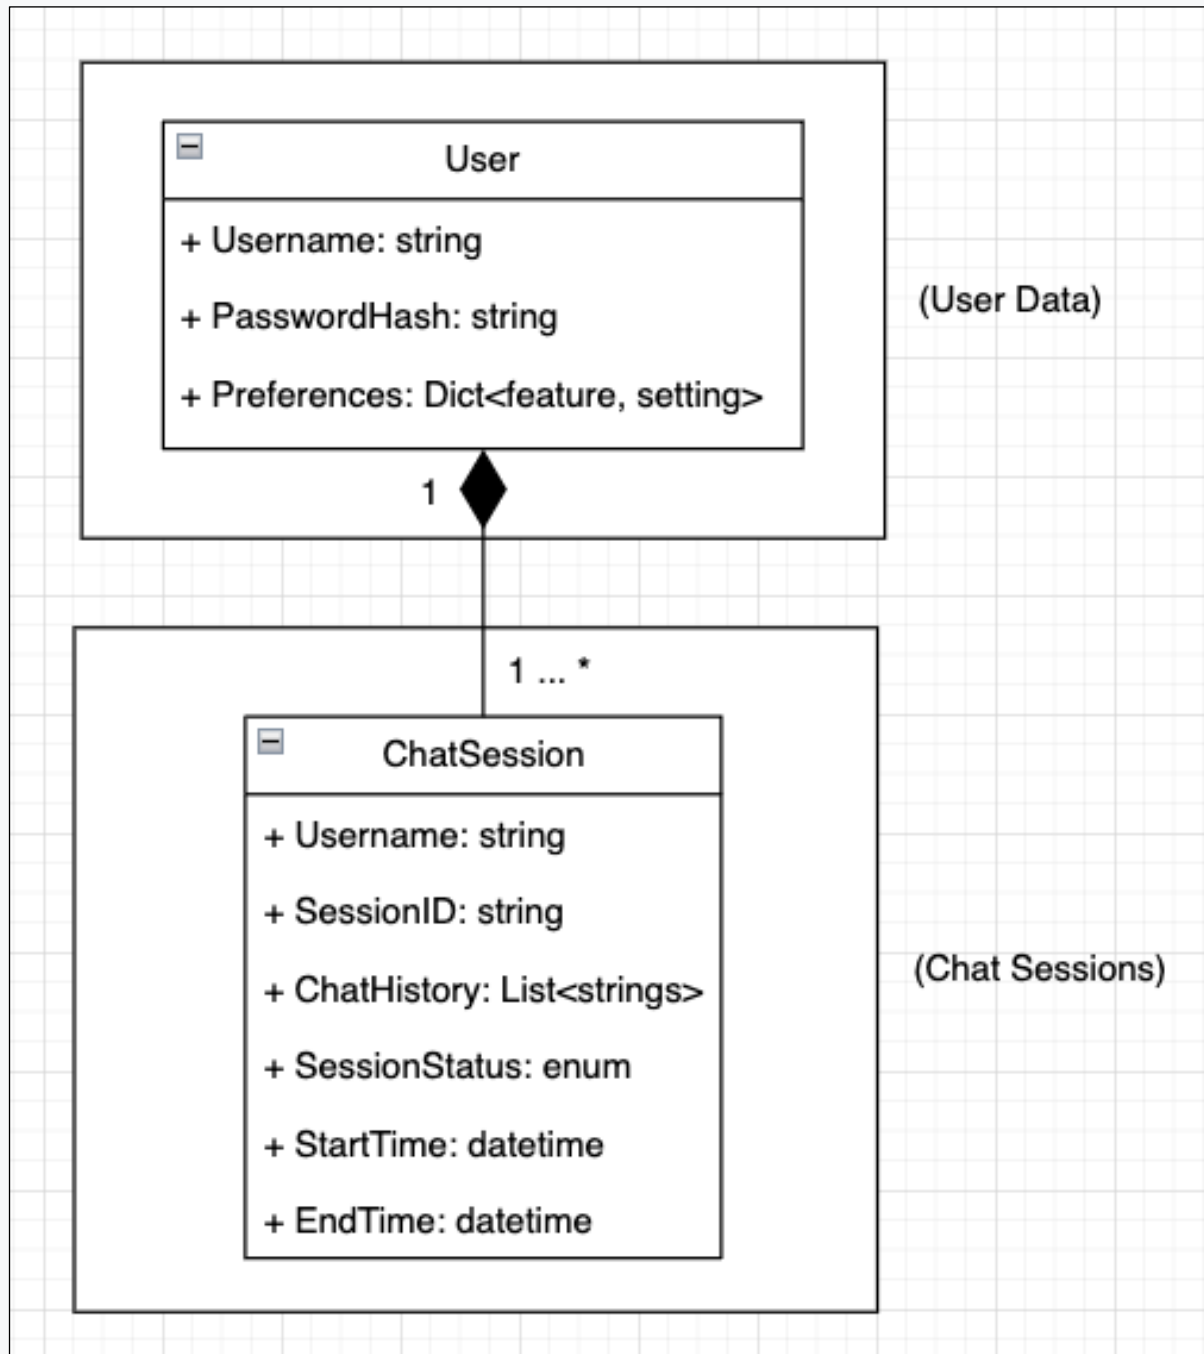
Documentation of System Architecture, Database, UI, & Testing



System Architecture

Technologies chosen:

- React for UI
- BERT for embedding
- spaCy and NLTK for NLP
- DBpedia for KG
- FAISS for VS
- LLaMA for LLM



Database Design

Optimization strategies:

- Indexing on Username and SessionID

User Interface Design

The image shows a chat interface mockup. It features a vertical list of messages. Each message consists of a blue bubble on the left and a gray bubble on the right. The blue bubbles contain user input, and the gray bubbles contain the AI's response. The messages are as follows:

- Blue bubble: "Hello, my name is Chandler!"
- Gray bubble: "RAG Model gives appropriate response"
- Blue bubble: "Who is Chandler"
- Gray bubble: "RAG Model gives appropriate response"
- Blue bubble: "What is the meaning of life?"
- Gray bubble: "RAG Model gives appropriate response"
- Blue bubble: "What is love?"
- Gray bubble: "RAG Model gives appropriate response"

At the bottom of the interface is a white input field with the placeholder text "Type your message..." and a blue "Send" button to its right. A vertical scrollbar is visible on the right side of the message list.

Testing Plan

Unit

- Testing isolated handlers
- Covering edge cases and boundary conditions

Integration

- Grouping interconnected components to test data flow
- Mocking outside components and dependencies

System

- Manual functional testing
- Measuring response time
- Stress testing
- Gathering feedback from stakeholders and end users

Tools:

- *Unittest* or *PyTest* as a testing framework
- *unittest.mock* for mocking dependencies
- *Coverage.py* to measure code coverage
- *spaCy* and *NLTK* to compare query responses
- *GitHub Actions* for CI/CD
- *Jmeter* or *Locust* to assess response time and scalability

Sprint Achievements & Challenges

- **Looking back on Sprint 2.**
- **Solution Approach Section**
- **Testing Approach Section**
- **Extracting Wikipedia dataset**
- **NLP processing of a query**
- **Generating text embeddings for sample of dataset**
- **DBpedia querying**
- **Basic frontend prototype of RAG App**

Next Steps & Retrospectives

- **Upcoming Tasks:**
- **Sprint Retrospective**
- **Client Feedback Consideration**
- **Generate text embeddings**
- **Integrate vector search**
- **Integrate knowledge graph**
- **LLM processing for responses**

Conclusion

- **Sprint 2 Achievements:**
- **Next Steps:**
 - **Optimize Embedding Generation:** Scale up for the full Wikipedia dataset with GPU acceleration.
 - **Implement FAISS Indexing:** Enable fast vector search for relevant results.
 - **Link to Knowledge Graph:** Enhance response accuracy by integrating vector search with the knowledge graph.
 - **Refine LLM Processing:** Improve query responses by combining vector and knowledge graph data.
 - **Advance Frontend:** Continue developing the chat interface for a smoother user experience.

Demo
